

Profile-Based Adaptive XAI for Cognitive Accessibility — Research Approach

Sep 29, 2026 · @Minidu

Summary (version 3, 4 Oct 2026)

We will write an **8-page full paper** for UISE 2027: a requirements-driven way to adapt the explanation of a flood ML prediction to a user-chosen explanation profile. It is built as a React Native prototype and evaluated with a persona-based cognitive walkthrough.

What changed from version 2

- Task 1 triage is now Minindu's manual evaluation of all 128 guidelines, not a first-pass proposal.
- The guideline tables moved to a separate evidence catalogue, with how each guideline is adapted in our flood EWS app and the published evidence for it.
- 19 guidelines from the manual review of the catalogue were added to the 31 from the first triage (50 in total); none was removed.

What changed from version 1

- **Paper type:** 5-page vision paper → 8-page full paper, because we now build and evaluate a prototype.
- **Prediction types are grounded:** chosen from the literature and from deployed AI flood systems (Google Flood Hub), not picked by us.
- **New motivation study:** we run SHAP, LIME and a counterfactual on a synthetic flood model and show which of our guidelines their output fails.
- **Profile questions are grounded:** built from the cognitive domains in standard screening tools and from validated self-report scales, not from the tests themselves (see Task 3 for why).
- **Three profiles are justified** from existing apps (Competitive Analysis 04) and prior XAI work.
- **Prototype:** two functions (set up / select profile; view prediction), wireframes first, then AI-assisted build on our existing React Native EWS app, same colours as v2.
- **Evaluation:** cognitive walkthrough with personas built from the three profiles. Evaluators record problems in plain words; we map them to guidelines afterwards.

Unchanged: keep the uncertainty ("Possible", not deleted), profiles by need not diagnosis, fidelity invariants as the core SE idea, Brief 2 (Vedda) stays a separate paper.

The schedule and owners are in the separate plan doc, "Adaptive XAI Paper — Feasible Plan to 13 Nov".

Research question and scope

Central question (Brief 1): How can cognitive accessibility requirements be identified, represented as XAI preference profiles, and used to adapt the explanation of disaster-related ML predictions?

Sub-questions:

1. Which existing cognitive accessibility requirements are relevant to XAI?
2. Which disaster ML prediction types create different explanation requirements?

3. How can XAI profiles determine explanation content, detail and interaction?

Motivating example. A standard output reads: "In two days there is a 30% chance your area will flood due to current weather." For someone with a cognitive impairment this may be hard to process: a percentage, a relative time, an implied location and no action. The adapted version keeps every fact but changes order, form and depth. A worked example is under Task 3.

In scope

- People with cognitive impairments (dementia, MCI, autism, ADHD and others), described by explanation *needs*, not diagnosis.
- One or two flood prediction types, with synthetic values clearly labelled.
- A literature-grounded requirements synthesis, profile specification, Figma demonstration and framework.

Out of scope for this paper

- Changing the ML model or its prediction.
- Claims of validation with users (no study has been run).
- Indigenous communities such as the Vedda. Their needs (source, provenance, time period, community governance) are a different problem, covered by Research Brief 2 as a separate paper.

Positioning and the gap

XAI for people with cognitive disabilities already exists, so our novelty has to be the combination of safety-critical predictions, a requirements process and fidelity checking.

Work	What it does	What it leaves open for us
NeuroAdaptX (Bunde, DESRIST 2026)	Adapts explanation structure, modality and density to self-reported profiles (ADHD, autism, dyslexia); RCT, N=216; keeps explanation content the same	Closest competitor. No safety-critical or disaster setting, no RE process, adapts presentation only. Must be cited and contrasted.
Tielman et al., "Explainable AI for all" (Technology in Society, 2024)	Roadmap for inclusive XAI for people with cognitive disabilities; recommends personalised, adaptive explanations	No disaster domain, no requirements specification or profile structure
ProfileXAI (Corrales et al., 2025)	SHAP/LIME/Anchor + LLM explanations for three user types (ML engineer, domain expert, lay user)	Profiles by expertise, not accessibility; healthcare
Accessible XAI survey (arXiv 2407.17484)	Surveys accessibility of XAI, mostly visual impairment	Cognitive accessibility under-covered
Madugalla et al. (2025), human-centred RE for disaster EWS	RE process for disaster EWS apps	Explicitly leaves cognitive impairment out; no XAI
Our earlier paper	RE process + prototype for cognitively accessible EWS	Warns and guides evacuation, but never explains a prediction

Gap statement for the paper: no work turns cognitive accessibility requirements into explanation requirements for safety-critical ML predictions, specifies them as user-chosen profiles, and checks that adapted explanations stay faithful to the prediction.

Differentiation from NeuroAdaptX in one line: they show presentation adaptation helps comprehension; we specify *what* must be explained, what software capability each explanation need depends on, and what can never be dropped in a safety-critical warning.

What our earlier work already shows

Across our paper, interim report, four competitive analyses, 128-guideline catalog and prototype, nothing explains a prediction. This is the new gap, and our own material is the evidence for it.

Source	Relevant finding	Use in the new paper
Final paper ("Designing Inclusive Disaster EWS...")	128-item catalog from 969 guidelines / 49 sources; rule-based profiles; AI only for optional plain-language rewriting	Starting point for Task 1; the prior-work paragraph
Competitive Analysis 01	5 production EWS apps (FEMA, MyShake, Google Public Alerts, Red Cross, Disaster Alert): none cognitively accessible; plain-language threat descriptions partial or absent	Evidence that mainstream EWS output is not written for this population
Competitive Analysis 02	13 apps split into predictive EWS vs cognitive-accessibility apps; none does both	Extend the argument: even the predictive apps never explain their predictions
Competitive Analysis 03	Four-layer architecture (baseline → profile rules → alert mode → optional AI); "accessibility illusion" risk of LLM simplification; only 6% of AI-in-EWS work targets dissemination	The architecture the new framework plugs into; motivation for fidelity invariants
Competitive Analysis 04	UI techniques across 7 apps; MyShake shows magnitude colours, Disaster Alert map symbols, no "why" or "how sure" anywhere	Evidence for the explanation gap
128-guideline catalog	Covers how to present (plain language, numbers in several forms, progressive disclosure), not what to explain	Task 1 triage (next section)
Prototype v2 (Figma)	Onboarding picks a condition (Dementia, Autism, MCI, ADHD, Schizophrenia); alert states Safe / Warning / Danger; danger screen says "Danger! Flood coming!"	Base for Task 4; needs a forecast state and an explanation-profile screen

The Figma file itself could not be opened by automated tools. Prototype observations come from Figure 3.2 of the interim report and Figure 2 of the final paper.

Task 1: Which requirements matter for XAI

Minindu manually evaluated all 128 guidelines and found 50 that affect how a prediction is explained: 31 from the first triage (10 XAI-related, 21 both) and 19 more from a second manual review of the colour-coded catalogue. The rest are general UI.

The full list is in a separate document, the [XAI guideline evidence catalogue](#) (also exported as a PDF). For each guideline it gives the implication for the explanation, how to adapt it in our EWS app, and published evidence: 38 have direct evidence, 9 indirect, and 3 have none found and stay as our inference.

Classes follow Brief 1: **A** = general UI, **B** = XAI-related (changes what is explained or how it is understood), **C** = both. None of the 49 source papers discusses AI explanations, so every B and C implication is *inferred by us*. The paper must say this.

Why check the evidence again. The 128 guidelines are already sourced, but their sources answer a different question: does this help people with cognitive impairments use an interface or follow a warning? We apply them to a new object, the explanation of an ML prediction, and that transfer is our inference. The evidence catalogue tests the transfer: does any study show the property changes how people understand an AI explanation, a probability or a warning? Direct evidence supports the transfer, indirect evidence supports it only partly, and "none found" marks it as inference.

What the catalog does not cover

The catalog explains *how* to present information, never *what* to explain. A targeted literature search is needed for:

- **Reasons and evidence:** why the risk is high, which observed conditions support it. Feature importance must not be presented as cause.
- **Uncertainty for low numeracy:** icon arrays, natural frequencies, verbal probability words and how they are misread.
- **Forecast change over time:** explaining upgrades and downgrades.
- **Source and provenance:** who issued the forecast and from what data.
- **False alarms and trust:** explaining a flood that did not happen so the next warning is still heeded.
- **Impact-based framing:** what the flood will *do* to the user, not what the weather will *be* (WMO impact-based forecasting).
- **Fidelity of simplified text:** checking an Easy Read or LLM version keeps the facts.

Task 2: Prediction types, grounded in literature and deployed AI

We choose the two output types that the leading deployed AI flood system actually produces: a **severity class** and a **multi-day river forecast with uncertainty**. That makes the choice defensible rather than ours.

Evidence from deployed AI systems

System	Model	What it outputs	Type it supports
Google Flood Hub (Nearing et al., Nature 2024)	LSTM encoder–decoder; inputs ECMWF weather, satellite rain, basin attributes	Daily streamflow up to 7 days, as a probability distribution; events judged against 1-, 2-, 5-, 10-year return periods; >80 countries; CAP alerts and push notifications	Time-series forecast with uncertainty
Flood Hub API (as documented by a third-party client)	Same	Gauge severity: NO_FLOODING / ABOVE_NORMAL / SEVERE / EXTREME; trend; 8-day forecast; inundation polygons graded HIGH / MEDIUM / LOW	Risk classification; spatial prediction
MyShake (Competitive Analysis 03)	Neural network on phone accelerometers	Earthquake vs everyday motion, then alert	Event classification
Sri Lanka Irrigation Department river gauges	Threshold rules (not ML)	Alert / minor flood / major flood levels	The local warning vocabulary users already hear (to verify)

Evidence from literature

- Nearing et al. (2024) above: the main citation; shows AI forecasts reaching the public with uncertainty available.
- Mosavi, Ozturk & Chau (2018), "Flood prediction using machine learning models: literature review", *Water*: classification and time-series regression as the dominant ML flood output types (**from memory; Imansha to confirm**).
- Recent XAI flood studies apply SHAP to flood susceptibility models, e.g. [this 2025 study](#), which shows feature attribution is the standard explanation in this domain.

Choice and rationale

Prediction type	Chosen?	Why
Risk classification	Yes	Flood Hub's severity classes; maps directly to CAP severity and local warning levels
Time-series forecast	Yes	Flood Hub's core output; carries uncertainty and change between forecasts
Spatial prediction	Future work	Needs maps, which are hard for this population; adds a third explanation problem
Event likelihood (landslide)	Out of scope	Different hazard; keep one hazard for the paper

Scenario (synthetic, labelled in the app): a Flood Hub–style forecast for one Sri Lankan river gauge: a 30% chance of reaching SEVERE (exceeding the danger level) in two days, upgraded from ABOVE_NORMAL

yesterday.

Task 2b: What standard XAI gives vs what our guidelines need

This study motivates the framework: we run the standard XAI methods on the same flood prediction and show which explanation requirements their output does not meet.

Method

1. Train a small model on a synthetic or public flood dataset (inputs such as 72-hour rainfall, upstream level, soil moisture; output = flood class + probability).
2. Explain one prediction (the 30% case) with **SHAP** (Lundberg & Lee, 2017), **LIME** (Ribeiro et al., 2016), a **counterfactual** (Wachter et al., 2017) and optionally **Anchors** (Ribeiro et al., 2018).
3. Record each method's raw output exactly as produced, as a paper figure.
4. Score each output against the B-class guidelines and the five invariants.

Expected shape of the result (to be replaced with real outputs):

Need (guideline)	SHAP	LIME	Counterfactual	Our adapted explanation
Bottom line first (G21)	No: list of feature weights	No	Partly	Yes
Likelihood without numeracy (G26, G119)	No: raw probability, log-odds	No	No	"Possible" + 3 of 10 picture
Place and time window (G73)	No	No	No	"Here, Thursday"
Consequence and action (G50)	No	No	No	"Get your go-bag ready"
Familiar references (G38)	No: feature names and units	No	Partly ("if rain < 80 mm")	River vs bridge mark
What changed (G89)	No: single prediction only	No	No	"Yesterday this was lower"
Why, in plain words (G25)	Feature attribution, not cause	Local weights	Closest: a "what if"	Reasons from observations

The point for the paper: standard XAI answers "which input features moved the model", for developers. Our guidelines need "what does this mean for me, how sure, and what do I do". That gap is why a requirements-driven adaptation layer is needed.

Task 3: XAI profiles, grounded in literature

The standard cognitive tests cannot be used as the profile itself. Their cognitive domains, and validated self-report scales, give each profile question its literature backing.

Can we use the Alzheimer's Association screening tools?

The [Alzheimer's Association list](#) covers patient tests (Mini-Cog, GPCOG part 1), informant tools (AD8, GPCOG part 2, Short IQCODE) and FDA-cleared digital tests. Using them directly fails on five counts:

Problem	Evidence
They are diagnostic screens given by clinicians, not app setup	Mini-Cog: "healthcare provider after brief training"; results decide whether a full diagnosis is needed (Alzheimer's Association)
Not valid for people with intellectual disability	General dementia tests cannot separate dementia from pre-existing disability; ID-specific tools such as DSQIID are needed (Zeilinger et al., 2022)
Biased by education and language	MMSE should not be used clinically without about grade-8 education and English fluency (MMSE overview)
Licensing	MMSE is copyrighted and sold through PAR (same source)
A score measures deficit, not explanation preference	Brief 1 forbids assigning profiles from diagnosis; catalog G84 forbids labelling by condition

What we use instead:

1. **The cognitive domains these tests measure** (memory, attention, executive function, orientation to time and place, language, numeracy). Each domain gives one profile question, so every question traces to a recognised domain.
2. **Validated self-report preference scales** that ask about preferences without testing ability. The model is the [Subjective Numeracy Scale](#) ([Fagerlin et al., 2007](#)), "measuring numeracy without a math test", which includes items on preferring numbers vs words.
3. **The informant idea** (AD8, IQCODE): a carer can answer or help, which matches G107.
4. **Function, not diagnosis:** WHO's ICF describes disability by functioning, which supports need-based profiles ([Supun to add the citation](#)).

The question template

Every question follows one pattern, so one well-made question becomes the model for the rest:

Domain → explanation decision it controls → Easy Read question with picture options → "Not sure, choose for me" default → "Ask my helper" option

Worked question (numeracy → likelihood form). Grounded in the Subjective Numeracy Scale's preference items, G26 and G119.

- Question: "When we tell you how likely a flood is, what helps you most?"
- Option A: words ("Possible", "Likely") — picture of the word card
- Option B: words and a picture (3 of 10 houses)
- Option C: words, picture and the number (30%)
- Default if unsure: B. The user can change it later.

All questions from the template

Each question now traces to the guidelines it serves and to the published evidence behind them, from the evidence catalogue. **D** = direct evidence, **I** = indirect evidence; *new* = added in the second manual review.

Domain (screening-test source)	Explanation decision	Question (draft)	Guidelines (evidence)	Literature backing
Numeracy (SNS)	Likelihood form	When we tell you how likely a flood is, what helps you most?	G26 (D), G119 (D), G81 (D, <i>new</i>)	Icon arrays help low-numeracy and older adults (Galesic et al., 2009); probability words are read inconsistently (Budescu et al., 2009); "30% chance" is read in contradictory ways (Gigerenzer et al., 2005)
Attention (Mini-Cog, MoCA attention items)	Detail on first view	Do you want only the main message first, or everything at once?	G41 (D), G21 (D), G11 (D), G18 (D, <i>new</i>)	Less, focused information improves comprehension, most for low numeracy (Peters et al., 2007); complex explanations are harder to interpret (Lage et al., 2019)
Executive function / planning (Mini-Cog clock drawing)	Action detail	Do you want us to tell you exactly what to do, step by step?	G4 (D), G50 (D), G22 (D, <i>new</i>)	Explanations explored step by step were understood better (Mindlin et al., 2024); protective action guidance improves understanding (Sutton et al., 2021); list instructions are recalled better (Morrow et al., 1998)
Memory (Mini-Cog recall; AD8 "repeats questions")	Reminders and change alerts	Should we remind you of the warning and tell you if it changes?	G109 (I), G89 (D), G54 (D), G27 (I, <i>new</i>)	Earlier forecasts anchor judgments of new ones (Herdener et al., 2018); repetition helps people with dementia learn (Creighton et al., 2013); too many alerts get ignored (Ha et al., 2026)
Orientation to time and place (MMSE orientation)	Time and place wording	How should we say when? "Thursday morning" or "in 2 days"?	G47 (D), G73 (D), G114 (D)	Orientation and calculation are among the most impaired items in MCI and dementia (Liu et al., 2019); location in a warning changes risk judgments (Klockow-McClain et al., 2020); time indications improve understanding (Dallo et al., 2022)
Language	Modality and reading level	Do you like to read, listen, or see pictures?	G128 (D), G67 (D), G14 (D, <i>new</i>), G15 (I, <i>new</i>), G17 (I, <i>new</i>)	Text + visual explanations help lay users most (Szymanski et al., 2021); matching the preferred format improves understanding (Tait et al., 2012); jargon impairs processing (Bullock et al., 2019)

Domain (screening-test source)	Explanation decision	Question (draft)	Guidelines (evidence)	Literature backing
Reasoning / curiosity	Explanation goal	Do you want to know why the app thinks a flood may come?	Brief 1; G25 (D)	Adaptive XAI is the key to inclusive XAI (Tielman et al., 2024); lay users cannot be expected to read a causal chain (Miller, 2019); impact information improves warning understanding (Weyrich et al., 2018)

Guidelines without published backing (3)

Three of the 50 XAI-relevant guidelines were marked as possibly relevant, but the evidence search found no source linking them to understanding an explanation, forecast or warning. They stay in the requirements as **our inference**, and the paper must label them that way.

ID	Guideline	Why it was marked	Where this document relies on it	How we treat it
G20	Include symbols and letters necessary to decipher	Sinhala and Tamil explanations must render correctly, including numerals, % and dates	Supports G121 (local-language Easy Read)	An engineering requirement: test that every explanation string renders in all three languages. Closest source is context only: CAP alerts lost integrity on small handsets in Sri Lankan trials (Waidyanatha et al., 2007), with no comprehension measure.
G42	Info to prepare for a task	State the expected timing and when the next forecast update arrives	Time-window invariant (with G73); "Next update: tonight at 6 pm" in the "How sure, and what changed?" profile	The invariant stays grounded through G73, which has direct evidence. Next step: a targeted search on lead-time and update-time communication.
G65	Focus Not Obscured (WCAG 2.4.11)	"Why?" and "How sure?" must never be hidden under a banner	Layout rule only	Treat as general UI (class A) in the paper; keep it in the prototype as a WCAG check.

Three profiles, and why these three

The profiles are bundles of default answers that the user can still change. Each bundle mirrors a pattern seen in existing apps or prior XAI work.

Profile	Answers it bundles	Why this profile exists
"Just tell me what to do"	Main message first; step-by-step actions; words only; listen	One-action-per-screen apps (Emergency Chat, MapHabit, GoTDAH) and calm "overwhelm mode" in Unmasked (Competitive Analysis 04); G124

Profile	Answers it bundles	Why this profile exists
"Show me why"	Reasons on; words + picture; landmark comparisons	Photo/visual-step design in MapHabit and Dementia Emergency (CA04); the "why" goal in Tielman et al. 2024 and NeuroAdaptX
"How sure, and what changed?"	Words + picture + number; change alerts on	Countdown/time externalisation in GoTDAH, data visualisation in Unmasked (CA04); uncertainty needs in the XAI literature

In the paper these profiles are presented as **hypothetical**, derived from literature and apps, not from users.

Worked example: one 30% prediction, three explanations

Same synthetic prediction everywhere: moderate risk, 30% likelihood of flooding in this area in two days, upgraded from yesterday's "low".

Standard output: "In two days there is a 30% chance your area will be flooded due to the current weather."

"Just tell me what to do"

1. "Flooding is possible here on Thursday. It is not certain."
2. "Get your go-bag ready today."
3. Button: "Why?" (opens the reasons)

"Show me why"

1. "Flooding is possible here on Thursday."
2. "Why: a lot of rain has fallen upstream, and the river is rising."
3. River-level picture against the bridge marker (synthetic).
4. Button: "How sure is this?"

"How sure, and what changed?"

1. "Flooding is possible here on Thursday. 3 in 10 chance." (icon array: 3 of 10 houses filled)
2. "Yesterday this was low. It went up because of new rain."
3. "Next update: tonight at 6 pm."
4. Exact figure on tap: 30%.

All three keep the hazard, the place, the time, the likelihood and the action. Only order, form and depth change.

Why "Possible": as we recall, CAP 1.2 defines *Possible* as possible but not likely ($p \leq \sim 50\%$) and *Likely* as $p > \sim 50\%$. That would make 30% → "Possible" a standard mapping. **To verify against the OASIS CAP 1.2 specification before citing.**

Task 4: Prototype (React Native) from wireframes

We extend our existing React Native EWS app with two functions. Every screen is wireframed first, because AI-generated builds from text prompts alone came out generic in our earlier work.

Function 1: Set up and select an explanation profile

- First-time setup: the template questions, one per screen, with picture options and "Not sure, choose for me".
- Pick one of the three ready profiles instead of answering questions.
- Change the profile later from settings; carer-assisted setup (optional).

Function 2: View a prediction and its explanation

- A new "Forecast" state before Warning, holding "Flooding is possible here on Thursday".
- The same synthetic prediction rendered per profile; "Why?" and "How sure?" on request.
- When it becomes a live alert, every profile collapses to action only (G124).
- A "Synthetic data" label on every forecast screen.

Build method

1. Wireframe every screen (hand sketches or a tool such as Excalidraw or Figma).
2. Write a short spec per screen: purpose, components, text, which profile attribute it reads.
3. Prompt the AI with the wireframe image + spec + the v2 colour tokens and component list; one screen per prompt.
4. Review each generated screen against its wireframe before merging.

Task 4b: Cognitive walkthrough

We evaluate the prototype with a persona-based cognitive walkthrough, adapting the method from our earlier paper, [Wharton et al. \(1994\)](#), [IxDF guidance](#), [Madugalla et al. \(2025\)](#) and the persona-based walkthrough in [Huynh et al. \(COMPSAC 2021\)](#).

Tasks (defined from the two functions, before the walkthrough):

ID	Task	Function
T1	Set up an explanation profile for the first time	1
T2	Switch to a different ready-made profile	1
T3	Read a new forecast and know what, where, when, how likely, and what to do	2
T4	Ask for more explanation ("Why?" or "How sure?") and return	2
T5	Notice that the forecast has changed since yesterday	2

Personas: one per profile, built like our earlier evidence-record personas (affected domains, documented behaviours, urgency modifier, sources), plus an optional carer persona.

Questions per step: the four standard questions plus our fifth (detect and recover from an error). For T3 we add a comprehension check: can the persona state the five invariants after reading?

Evaluator sheet (simpler than last time):

Column	What the evaluator writes
Task, step, persona	Pre-filled

Column	What the evaluator writes
Q1–Q5	Yes / No
What went wrong	Plain words
Severity	0–4
Screenshot	Required for severity 3–4


No guideline column. Evaluators only describe problems; the team maps each defect to a guideline afterwards, in a separate pass.

Bias control: three members walk independently; the person who built a screen does not lead it; severities are rated before discussion, and agreement is reported.

Task 5: The framework

The framework turns accessibility requirements into explanation behaviour, and checks every adapted explanation against fixed invariants before the user sees it.

Profiles change the explanation; invariants keep the facts



Proposed framework · Brief 1 chain plus capability dependencies and an invariants check

Requirements flow down to a profile; each explanation need also creates a capability the system must have. Profile and prediction meet in the adaptation rules. Verification findings feed back into the requirements.

Fidelity invariants

No profile, simplification or AI rewording may remove these.

Invariant	Example in the scenario	Grounding
Hazard	Flood	G11, G119
Place	"Here" / your area	G73
Time window	Thursday	G73, G42
Likelihood (uncertainty kept)	"Possible", with picture or number per profile	G26, G119
Action	Get your go-bag ready	G11, G50

Requirement → capability dependencies

This is the SE contribution: an explanation need is not just a UI choice; it requires data and functions in the system.

Explanation need	Capability the system needs
What changed since last time	Stored previous forecasts and a comparison
Why is the risk high	Model exposes supporting observations (not feature importance as cause)
How sure	Probability or range from the model
Level against a landmark	Local reference levels for each area
When is the next update	Forecast schedule metadata
Carer explanation	Consent record and a second profile

Fit with our four-layer architecture

- **Layer 1, baseline:** one fixed accessible explanation containing all invariants; used if everything above fails.
- **Layer 2, profile rules:** explanation adaptation lives here. It is deterministic and auditable.
- **Layer 3, alert mode:** when an alert is live, every profile collapses to action only (G124).
- **Layer 4, optional AI:** LLM wording is shown only if it passes the invariants check; otherwise the Layer 2 text is shown.

Worked trace ("How sure, and what changed?" profile)

1. Requirement G26: provide alternatives for numerical concepts.
2. Explanation need: understand likelihood without relying on numeracy.
3. Profile attribute: likelihood form = word + picture + number.
4. Prediction: p = 0.30 (synthetic), upgraded from "low".
5. Rule: CAP word "Possible" + icon array 3 of 10 + "30%" on tap + change note.
6. Invariants check: hazard, place, time, likelihood, action present → pass.
7. Screen shown; the requirement traces to the screen element.

Other research directions (for the agenda section)

A vision paper can list these as future work; each could become its own study.

Direction	Question	Link to our work
Explanation under stress	How much explanation is possible once an alert is live?	Extends the alert-state layer (G124, G125)
Lead time	Should a 30% forecast be explained differently at 48 h, 12 h and 1 h?	New; time-series scenario

Direction	Question	Link to our work
False alarms and trust	How to explain a flood that did not happen so the next warning is still heeded?	Appropriate reliance, Brief 1 long-term vision
Caregivers	One prediction, two explanations: person and helper	G80, G107
Profile privacy	A saved profile may reveal a disability; who can see it?	UISE lists "privacy and ethics in adaptive UIs"
LLM explanation fidelity	Can generated explanations be checked automatically against the prediction and invariants?	Layer 4 of the architecture; the "accessibility illusion"
Language and speech	Sinhala and Tamil Easy Read explanations, read aloud	G121, G108
Evaluation	Requirement satisfaction, fidelity, comprehension, appropriate reliance, decision support	Brief 1 long-term vision
Indigenous communities	Community-defined explanation needs (source, provenance)	Brief 2, separate paper

Paper plan (8-page full paper)

Submission is due 13 Nov 2026 (AoE). The full-paper limit is 8 pages including references. Week-by-week work and owners are in the separate plan doc.

Section	Pages	Content
Introduction	0.75	30% example, problem, contributions
Related work	1	XAI for cognitive disability (NeuroAdaptX, Tielman), disaster EWS RE, our earlier paper
Motivating study: standard XAI vs guidelines	1	SHAP / LIME / counterfactual outputs, comparison table
Requirements and profiles	1.25	Task 1 triage summary, prediction-type choice, question template, three profiles
Framework	1	Framework figure, invariants, requirement → capability table
Prototype	1	Screens per profile, build method
Cognitive walkthrough	1.25	Tasks, personas, defects, agreement
Discussion, limitations, future work	0.5	No user validation yet; agenda
References	0.75	About 30

Other dates: notification 11 Dec 2026, camera-ready 29 Jan 2027.

Rules to respect and open questions

Rules from the briefs (reviewers will check these)

- Never claim profiles or explanations were validated with users; no study has been run.
- Adapt the explanation, never the prediction or its risk estimate.
- Don't assign profiles from a diagnosis.
- Mark every requirement as stated by a source or inferred by us; label hypothetical profiles as hypothetical.
- Label synthetic forecast values; don't invent evidence or data sources.
- Don't present feature importance as cause.

Risks

- **Reviewer knows NeuroAdaptX:** without a clear contrast the paper reads as incremental. Mitigation: the positioning table and the one-line differentiation.
- **5-page limit:** four tables plus two figures will not fit. Keep one summary table from Task 1 and move the full catalog triage to an online appendix.
- **CAP definition unverified:** check the OASIS CAP 1.2 spec before relying on the "Possible" mapping.

Open questions for the team

- ☐ Does Madam agree to drop the condition picker in favour of need-based profiles, or keep both?
- ☐ Which river basin and which documented forecast system do we base the synthetic scenario on?
- ☐ Does a second team member classify the 128 guidelines independently, so we can report agreement with Minindu's manual triage?
- ☐ Is the "fidelity invariants" idea the main SE contribution, or does Madam prefer the requirement-to-capability mapping (Brief 2's framing)?
- ☐ Where do we host the online appendix (catalog triage, profile spec)?

Sources

External

- [UISE 2027 call for papers](#)
- [Bunde \(2026\), NeuroAdaptX: neuro-adaptive explanations for cognitive accessibility in XAI interfaces, DESRIST](#)
- [Tielman et al. \(2024\), Explainable AI for all: a roadmap for inclusive XAI for people with cognitive disabilities, Technology in Society 79](#)
- [Corrales et al. \(2025\), ProfileXAI: user-adaptive explainable AI](#)
- [A survey of accessible explainable AI research](#)
- [OASIS Common Alerting Protocol v1.2, certainty values \(to be checked directly\)](#)

Our materials

- Research Brief 1 and Research Brief 2 (project docs)

- Final paper: "Designing Inclusive Disaster Early Warning Systems for People with Cognitive Impairment"
- Interim report (thesis 1)
- Competitive Analyses 01–04
- [EWS design guidelines, 128 categorised](#)
- [Design compendium](#)
- [Figma prototype](#)